

```
In [1]: 1 # If true, the WAV files will be read and their features will be save
        2 # As this is the most time consuming task, only enable it if you don'
```

```
In [2]: 1 # Defines the names of the CSV files
        2 TRAIN_CSV_FILE = "train.csv"
        3 TEST_CSV_FILE = "test.csv"
        4 MORE_TRAIN_CSV_FILE = "more_train.csv"
        5 MORE_TEST_CSV_FILE = "more_test.csv"
```

```
In [3]: 1 import matplotlib.pyplot as plt
        2 import numpy as np
        3 from matplotlib import cm
        4 import librosa
        5 import csv
        6 import os
        7
        8 def extractWavFeatures(soundFilesFolder, csvFileName):
        9     print("The features of the files in the folder "+soundFilesFolder
       10     header = 'filename chroma_stft rmse spectral_centroid spectral_ba
       11     for i in range(1, 21):
       12         header += f' mfcc{i}'
       13     header += ' label'
       14     header = header.split()
       15     print('CSV Header: ', header)
       16     file = open(csvFileName, 'w', newline='')
       17     writer = csv.writer(file)
       18     writer.writerow(header)
       19     genres = '1 2 3 4 5 6 7 8 9 0'.split()
       20     for filename in os.listdir(soundFilesFolder):
       21         number = f'{soundFilesFolder}/{filename}'
       22         y, sr = librosa.load(number, mono=True, duration=30)
       23         # remove leading and trailing silence
       24         y, index = librosa.effects.trim(y)
       25         chroma_stft = librosa.feature.chroma_stft(y=y, sr=sr)
       26         rmse = librosa.feature.rms(y=y)
       27         spec_cent = librosa.feature.spectral_centroid(y=y, sr=sr)
       28         spec_bw = librosa.feature.spectral_bandwidth(y=y, sr=sr)
       29         rolloff = librosa.feature.spectral_rolloff(y=y, sr=sr)
       30         zcr = librosa.feature.zero_crossing_rate(y)
       31         mfcc = librosa.feature.mfcc(y=y, sr=sr)
       32         to_append = f'{filename} {np.mean(chroma_stft)} {np.mean(rmse
       33         for e in mfcc:
       34             to_append += f' {np.mean(e)} '
       35         writer.writerow(to_append.split())
       36     file.close()
       37     print("End of extractWavFeatures")
       38
       39 if (CREATE_CSV_FILES == True):
       40     extractWavFeatures("./data/recordings/train", TRAIN_CSV_FILE)
       41     extractWavFeatures("./data/recordings/test", TEST_CSV_FILE)
       42     extractWavFeatures("./data/recordings/moreSpeakersTrain", MORE_TR
       43     extractWavFeatures("./data/recordings/moreSpeakersTest", MORE_TES
       44     print("CSV files are created")
       45 else:
```

```
46 print("CSV files creation is skipped")
```

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The text.latex.preview rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The mathtext.fallback\_to\_cm rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle: Support for setting the 'mathtext.fallback\_to\_cm' rcParam is deprecated since 3.3 and will be removed two minor releases later; use 'mathtext.fallback : 'cm' instead.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The validate\_bool\_maybe\_none function was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The savefig.jpeg\_quality rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The keymap.all\_axes rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The animation.avconv\_path rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

In C:\Users\Yana\Anaconda3\lib\site-packages\matplotlib\mpl-data\stylelib\\_classic\_test.mplstyle:  
The animation.avconv\_args rcparam was deprecated in Matplotlib 3.3 and will be removed two minor releases later.

The features of the files in the folder ./data/recordings/train will be saved to train.csv  
CSV Header: ['filename', 'chroma\_stft', 'rmse', 'spectral\_centroid', 'spectral\_bandwidth', 'rolloff', 'zero\_crossing\_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']  
End of extractWavFeatures  
The features of the files in the folder ./data/recordings/test will be saved to test.csv  
CSV Header: ['filename', 'chroma\_stft', 'rmse', 'spectral\_centroid', 'spectral\_bandwidth', 'rolloff', 'zero\_crossing\_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']  
End of extractWavFeatures  
The features of the files in the folder ./data/recordings/moreSpeakers Train will be saved to more\_train.csv  
CSV Header: ['filename', 'chroma\_stft', 'rmse', 'spectral\_centroid', 'spectral\_bandwidth', 'rolloff', 'zero\_crossing\_rate', 'mfcc1', 'mfcc2', 'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']

```

c17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
End of extractWavFeatures
The features of the files in the folder ./data/recordings/moreSpeakers
Test will be saved to more_test.csv
CSV Header: ['filename', 'chroma_stft', 'rmse', 'spectral_centroid',
'spectral_bandwidth', 'rolloff', 'zero_crossing_rate', 'mfcc1', 'mfcc2',
'mfcc3', 'mfcc4', 'mfcc5', 'mfcc6', 'mfcc7', 'mfcc8', 'mfcc9', 'mfcc10', 'mfcc11', 'mfcc12', 'mfcc13', 'mfcc14', 'mfcc15', 'mfcc16', 'mfcc17', 'mfcc18', 'mfcc19', 'mfcc20', 'label']
End of extractWavFeatures
CSV files are created

```

In [4]:

```

1 import pandas as pd
2 import csv
3 from sklearn import preprocessing
4
5 def preProcessData(csvFileName):
6     print(csvFileName+ " will be preprocessed")
7     data = pd.read_csv(csvFileName)
8     data['number'] = data['filename'].str[:1]
9     #Dropping unnecessary columns
10    data = data.drop(['filename'],axis=1)
11    data = data.drop(['label'],axis=1)
12    data = data.drop(['chroma_stft'],axis=1)
13    data.shape
14
15    print("Preprocessing is finished")
16    print(data.head())
17    return data
18
19 trainData = preProcessData(TRAIN_CSV_FILE)
20 testData = preProcessData(TEST_CSV_FILE)
21 moreTrainData = preProcessData(MORE_TRAIN_CSV_FILE)
22 moreTestData = preProcessData(MORE_TEST_CSV_FILE)

```

train.csv will be preprocessed

Preprocessing is finished

	rmse	spectral_centroid	spectral_bandwidth	rolloff	\
0	0.112672	741.829081	758.492178	1438.494873	
1	0.090344	635.610880	670.336296	1160.452403	
2	0.091456	667.786694	732.606545	1257.180176	
3	0.087751	712.304185	731.292437	1449.104818	
4	0.096603	844.363886	777.868127	1569.583263	

	zero_crossing_rate	mfcc1	mfcc2	mfcc3	mfcc4	\
0	0.034023	-295.578461	189.853683	-19.606564	6.078509	
1	0.033458	-339.148743	204.005249	-7.485528	14.297899	
2	0.033268	-327.507416	195.596924	-3.994768	21.315845	
3	0.035916	-320.809937	200.023743	-8.186146	12.661074	
4	0.049465	-315.801300	195.674118	-13.324564	3.544238	

	mfcc5	...	mfcc12	mfcc13	mfcc14	mfcc15	mfcc16	\
0	22.067095	...	-25.725817	-5.172223	-8.323026	-10.299589	-0.144793	
1	20.005120	...	22.100260	1.200001	5.515564	15.416007	0.405076	

## Section 2

There are 50 recordings for each digit for each speaker: Jackson, Nicolas and Theo (total 1500 recordings)

Training data has 49 recordings for each digit for each speaker: 1470 recordings total. Test data has 1 recordings for each digit for each speaker: 30 recordings total.

The data used here comes from the recordings stored in:

../data/recordings/train

../data/recordings/test

The model will be trained to predict the spoken digit.

```
In [5]: 1 # Splitting the dataset into training, validation and testing dataset
2 from sklearn.model_selection import train_test_split
3 X = np.array(trainData.iloc[:, :-1], dtype = float)
4 y = trainData.iloc[:, -1]
5 X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3)
6
7
8 X_test = np.array(testData.iloc[:, :-1], dtype = float)
9 y_test = testData.iloc[:, -1]
10
11 print("Y from training data:", y_train.shape)
12 print("Y from validation data:", y_val.shape)
```

```
Y from training data: (1029,)
Y from validation data: (441,)
Y from test data: (30,)
```

```
In [6]: 1 #Normalizing the dataset
2 from sklearn.preprocessing import StandardScaler
3 import numpy as np
4 scaler = StandardScaler()
5 X_train = scaler.fit_transform( X_train )
6 X_val = scaler.transform( X_val )
7 X_test = scaler.transform( X_test )
8
9 print("X from training data", X_train.shape)
10 print("X from validation data", X_val.shape)
11 print("X from test data", X_test.shape)
```

```
X from training data (1029, 25)
X from validation data (441, 25)
X from test data (30, 25)
```

```
In [7]: 1 #Creating a Model
2 from keras import models
3 from keras import layers
4 import keras
5
6 # model 1
7 model = models.Sequential()
```


```
8 model.add(layers.Dense(256, activation='relu', input_shape=(X_train.s
9 model.add(layers.Dropout(0.5))
10 model.add(layers.Dense(128, activation='relu'))
11 model.add(layers.Dropout(0.5))
12 model.add(layers.Dense(64, activation='relu'))
13 model.add(layers.Dropout(0.5))
14 model.add(layers.Dense(10, activation='softmax'))
15
16 # Learning Process of a model
17 model.compile(optimizer='adam',
18               loss='sparse_categorical_crossentropy',
19               metrics=['accuracy'])
20 model.summary()
21 #return model
22
23 # simple early stopping
24 from keras.callbacks import EarlyStopping
25
26 es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)
27
28 #Train with early stopping to avoid overfitting
29 history = model.fit(X_train,
30                     y_train,
31                     validation_data=(X_val, y_val),
32                     epochs=50,
33                     batch_size=128)
```

Using TensorFlow backend.

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	6656
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 10)	650

```
In [8]: 1 from matplotlib import pyplot
2 pyplot.plot(history.history['loss'], label='train')
3 pyplot.plot(history.history['val_loss'], label='test')
4 pyplot.legend()
```



```
In [9]: 1 def printPrediction(X_data, y_data):
2     print('\n# Generate predictions')
3     for i in range(len(y_data)):
4         prediction = model.predict_classes(X_data[i:i+1])
```

```
In [10]: 1 import numpy as np
2 from keras import backend as K
3 from keras.models import Sequential
4 from keras.layers.core import Dense, Dropout, Activation, Flatten
5 from keras.layers.convolutional import Convolution2D, MaxPooling2D
6 from keras.preprocessing.image import ImageDataGenerator
7 from sklearn.metrics import classification_report, confusion_matrix
8
9 def report(X_data, y_data):
10     #Confution Matrix and Classification Report
11     Y_pred = model.predict_classes(X_data)
12     y_test_num = y_data.astype(np.int64)
13     print('Confusion Matrix')
14     conf_mt = confusion_matrix(y_test_num, Y_pred)
15     print(conf_mt)
16     plt.matshow(conf_mt)
17     plt.show()
18
19     print('\nClassification Report')
20     target_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

```
In [11]: 1 print('\n# TEST DATA #\n')
2 score = model.evaluate(X_test, y_test)
3 print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))
4
5 # Prediction
```

```
# TEST DATA #
```

```
30/30 [=====] - 0s 2ms/step  
accuracy: 90.00%
```


```
# Generate predictions  
y=0, prediction=[0], match=True  
y=0, prediction=[3], match=False  
y=0, prediction=[0], match=True  
y=1, prediction=[1], match=True  
y=1, prediction=[4], match=False  
y=1, prediction=[1], match=True  
y=2, prediction=[2], match=True  
y=2, prediction=[2], match=True
```

```
In [12]: 1 print("Classification Report for Test Data\n")
```

```
Classification Report for Test Data
```

```
Confusion Matrix
```

```
[[2 0 0 1 0 0 0 0 0 0]  
 [0 2 0 0 1 0 0 0 0 0]  
 [0 0 2 1 0 0 0 0 0 0]  
 [0 0 0 3 0 0 0 0 0 0]  
 [0 0 0 0 3 0 0 0 0 0]  
 [0 0 0 0 0 3 0 0 0 0]  
 [0 0 0 0 0 0 3 0 0 0]  
 [0 0 0 0 0 0 0 3 0 0]  
 [0 0 0 0 0 0 0 0 3 0]  
 [0 0 0 0 0 0 0 0 0 3]]
```



Classification Report					
	precision	recall	f1-score	support	
0	1.00	0.67	0.80	3	
1	1.00	0.67	0.80	3	
2	1.00	0.67	0.80	3	
3	0.60	1.00	0.75	3	
4	0.75	1.00	0.86	3	
5	1.00	1.00	1.00	3	

## Section 3

There are 50 recordings for each digit for each speaker: Jackson, Nicolas and Theo (total 1500 recordings) Training data has 49 recordings for each digit for each speaker: 1470 recordings total. Test data has 1 recordings for each digit for each speaker: 30 recordings total.

In addition, there are 2 recordings for each digit for each speaker: Ankur, Caroline and Rodolfo (total 60 recordings) This addition training data has 1 recordings for each digit for each speaker: 30 recordings total. This addition test data has 1 recordings for each digit for each speaker: 30 recordings total.

Therefore the full data set has:

Training: 1500 recordings Training: 60 recordings

The data used here comes from the recordings stored in:

../data/recordings/train ../data/recordings/test ../data/recordings/moreSpeakersTrain ../data/recordings/moreSpeakersTest

```
In [13]: 1 # Splitting the dataset into training, validation and testing dataset
2 from sklearn.model_selection import train_test_split
3
4 fullTrainData = trainData.append(moreTrainData)
5
6 X = np.array(fullTrainData.iloc[:, :-1], dtype = float)
7 y = fullTrainData.iloc[:, -1]
8 X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3)
9
10 X_test = np.array(testData.iloc[:, :-1], dtype = float)
11 y_test = testData.iloc[:, -1]
12
13 X_more_test = np.array(moreTestData.iloc[:, :-1], dtype = float)
14 y_more_test = moreTestData.iloc[:, -1]
15
16 print("Y from training data:", y_train.shape)
17 print("Y from validation data:", y_val.shape)
18 print("Y from test data:", y_test.shape)
```



Y from training data: (1050,)

```
In [14]: 1 #Normalizing the dataset
2 from sklearn.preprocessing import StandardScaler
3 import numpy as np
4 scaler = StandardScaler()
5 X_train = scaler.fit_transform( X_train )
6 X_val = scaler.transform( X_val )
7 X_test = scaler.transform( X_test )
8 X_more_test = scaler.transform( X_more_test )
9
10 print("X from training data", X_train.shape)
11 print("X from validation data", X_val.shape)
12 print("X from test data", X_test.shape)
```

```
X from training data (1050, 25)
X from validation data (450, 25)
X from test data (30, 25)
X from other speakers test data (30, 25)
```


```
1 #Creating a Model
2 from keras import models
3 from keras import layers
4 import keras
5
6 # model 1
7 model = models.Sequential()
8 model.add(layers.Dense(256, activation='relu', input_shape=
9 (X_train.shape[1],)))
10 model.add(layers.Dropout(0.5))
11 model.add(layers.Dense(128, activation='relu'))
12 model.add(layers.Dropout(0.5))
13 model.add(layers.Dense(64, activation='relu'))
14 model.add(layers.Dropout(0.5))
15 model.add(layers.Dense(10, activation='softmax'))
16
17 # Learning Process of a model
18 model.compile(optimizer='adam',
19               loss='sparse_categorical_crossentropy',
20               metrics=['accuracy'])
21
22 # simple early stopping
23 from keras.callbacks import EarlyStopping
24 es = EarlyStopping(monitor='val_loss', mode='min', verbose=1)
25
26 #Train with early stopping to avoid overfitting
27 history = model.fit(X_train,
28                     y_train,
29                     validation_data=(X_val, y_val),
30                     epochs=50,
31                     batch_size=128,
```

```
In [15]: 1 # plot training history
2 from matplotlib import pyplot
```

```

3 pyplot.plot(history.history['loss'], label='train')
4 pyplot.plot(history.history['val_loss'], label='test')
5 pyplot.legend()

```



## Present the model performance

```

In [16]: 1 print('\n# TEST DATA #\n')
2 score = model.evaluate(X_test, y_test)
3 print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))
4
5 # Prediction
6 print('Prediction (X_test, y_test) = (50, 101) -> (50, 101)')

```

# TEST DATA #

30/30 [=====] - 0s 40us/step  
accuracy: 90.00%

```

# Generate predictions
y=0, prediction=[0], match=True
y=0, prediction=[3], match=False
y=0, prediction=[0], match=True
y=1, prediction=[1], match=True
y=1, prediction=[4], match=False
y=1, prediction=[1], match=True
y=2, prediction=[2], match=True
y=2, prediction=[2], match=True
y=2, prediction=[3], match=False
y=3, prediction=[3], match=True

```

```

In [17]: 1 print('\n# OTHER SPEAKERS DATA #\n')
2 score = model.evaluate(X_more_test, y_more_test)
3 print("%s: %.2f%%" % (model.metrics_names[1], score[1]*100))
4
5 # Prediction
6 print('Prediction (X_more_test, y_more_test) = (50, 101) -> (50, 101)')

```

```
# OTHER SPEAKERS DATA #
```

```
30/30 [=====] - 0s 95us/step
accuracy: 26.67%
```

```
# Generate predictions
y=0, prediction=[2], match=False
y=0, prediction=[2], match=False
y=0, prediction=[8], match=False
y=1, prediction=[0], match=False
y=1, prediction=[0], match=False
y=1, prediction=[1], match=True
y=2, prediction=[2], match=True
```


In [18]:

```
1 print("Classification Report for Test Data\n")
2 report(X_test, y_test)
3
4 print("Classification Report for Other Speakers\n")
```

Classification Report for Test Data

Confusion Matrix

```
[[2 0 0 1 0 0 0 0 0]
 [0 2 0 0 1 0 0 0 0]
 [0 0 2 1 0 0 0 0 0]
 [0 0 0 3 0 0 0 0 0]
 [0 0 0 0 3 0 0 0 0]
 [0 0 0 0 0 3 0 0 0]
 [0 0 0 0 0 0 3 0 0]
 [0 0 0 0 0 0 0 3 0]
 [0 0 0 0 0 0 0 0 3]]
```




## Classification Report

	precision	recall	f1-score	support
0	1.00	0.67	0.80	3
1	1.00	0.67	0.80	3
2	1.00	0.67	0.80	3
3	0.60	1.00	0.75	3
4	0.75	1.00	0.86	3
5	1.00	1.00	1.00	3
6	1.00	1.00	1.00	3
7	1.00	1.00	1.00	3
8	1.00	1.00	1.00	3
9	1.00	1.00	1.00	3
accuracy			0.90	30
macro avg	0.93	0.90	0.90	30
weighted avg	0.94	0.90	0.90	30

## Classification Report for Other Speakers

## Confusion Matrix

```
[[0 0 2 0 0 0 0 0 1 0]
 [2 1 0 0 0 0 0 0 0 0]
 [0 0 1 0 0 0 1 0 1 0]]
```



## Classification Report

	precision	recall	f1-score	support
0	0.00	0.00	0.00	3
1	1.00	0.33	0.50	3
2	0.17	0.33	0.22	3
3	1.00	0.33	0.50	3
4	0.00	0.00	0.00	3
5	0.33	0.33	0.33	3
6	0.20	0.67	0.31	3
7	0.00	0.00	0.00	3
8	0.20	0.33	0.25	3
9	1.00	0.33	0.50	3

accuracy			0.27	30
macro avg	0.39	0.27	0.26	30
weighted avg	0.39	0.27	0.26	30

```
C:\Users\Yana\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1437: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

In [ ]: